

Instructions

- Some parts of the code are already done for you
- You need to execute all the cells
- You need to add the code where ever you see "#### Add your code here ####"
- Marks are mentioned along with the cells

▼ Face detection

Task is to predict the boundaries(mask) around the face in a given image.

▼ Dataset

Faces in images marked with bounding boxes. Have around 500 images with around 1100 faces n

▼ Mount Google drive if you are using google colab

- We recommend using Google Colab as you can face memory issues and longer runtimes wh

```
from google.colab import drive  
drive.mount('/content/drive')
```

↳ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=9473

```
Enter your authorization code:  
.....  
Mounted at /content/drive
```

▼ Change current working directory to project folder (1 mark)

```
import os  
PATH = "/content/drive/My Drive/Colab Notebooks/"  
os.chdir(PATH)
```

▼ Load the "images.npy" file (2 marks)

- This file contains images with details of bounding boxes

```
# IMPORT LIBRARIES AND PACKAGES
```

```
import csv  
import math
```

<https://colab.research.google.com/drive/1JDija1tRxz4m62DzbuFSuwWRmdbX3VY#scrollTo=Z0Qa1UOue9TE&printMode=true>

```
import cv2

import numpy as np
import tensorflow as tf
from PIL import Image
from tensorflow.keras import Model
from tensorflow.keras.applications.mobilenet import MobileNet, preprocess_input
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.layers import Concatenate, Conv2D, UpSampling2D, Reshape
from tensorflow.keras.utils import Sequence
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import binary_crossentropy
data = np.load('./images_small.npy',allow_pickle=True)
```

- ▼ Check one sample from the loaded "images.npy" file (2 marks)

Hint - print data[10][1]

```
for key in data:
    print (key)
    print (data[10][1])
    break
```



```
[array([[ [42, 37, 34],
       [56, 51, 48],
       [71, 66, 63],
       ...,
       [23, 33, 34],
       [26, 36, 37],
       [28, 38, 39]],

      [[40, 35, 32],
       [51, 46, 43],
       [64, 59, 56],
       ...,
       [27, 36, 35],
       [24, 33, 32],
       [26, 35, 34]],

      [[43, 38, 35],
       [51, 46, 43],
       [61, 56, 53],
       ...,
       [28, 30, 27],
       [33, 35, 32],
       [35, 37, 34]],

      ...,

      [[56, 47, 40],
       [57, 48, 41],
       [61, 52, 45],
       ...,
       [67, 48, 42],
       [55, 35, 28],
       [60, 40, 33]],

      [[53, 44, 37],
       [54, 45, 38],
       [57, 48, 41],
       ...,
       [59, 40, 34],
       [60, 40, 33],
       [54, 34, 27]],

      [[53, 44, 37],
       [54, 45, 38],
       [57, 48, 41],
       ...,
       [59, 40, 34],
       [70, 50, 43],
       [64, 44, 37]]], dtype=uint8)
list([{'label': ['Face'], 'notes': '', 'points': [{'x': 0.08615384615384615, 'y': 0.10385756676557864}, {
```

▼ Set image dimensions (1 mark)

- Initialize image height, image width with value: 224

```
# SETTINGS

ALPHA = 1 # Width hyper parameter for MobileNet (0.25, 0.5, 0.75, 1.0). Higher width means

IMAGE_HEIGHT = 224
IMAGE_WIDTH = 224

HEIGHT_CELLS = 28
WIDTH_CELLS = 28

CELL_WIDTH = IMAGE_WIDTH / WIDTH_CELLS
CELL_HEIGHT = IMAGE_HEIGHT / HEIGHT_CELLS

EPOCHS = 1
BATCH_SIZE = 4
PATIENCE = 10

THREADS = 1

PATH = "/content/drive/My Drive/Colab Notebooks/"

TRAIN_CSV = PATH+"train.csv"
VALIDATION_CSV = PATH+"validation.csv"
```

▼ Create features and labels

- Here feature is the image
- The label is the mask
- Images will be stored in "X_train" array
- Masks will be stored in "masks" array

```
import cv2
from tensorflow.keras.applications.mobilenet import preprocess_input

masks = np.zeros((int(data.shape[0]), IMAGE_HEIGHT, IMAGE_WIDTH))
X_train = np.zeros((int(data.shape[0]), IMAGE_HEIGHT, IMAGE_WIDTH, 3))
for index in range(data.shape[0]):
    img = data[index][0]
    img = cv2.resize(img, dsize=(IMAGE_HEIGHT, IMAGE_WIDTH), interpolation=cv2.INTER_CUBIC
try:
    img = img[:, :, :3]
except:
    continue
X_train[index] = preprocess_input(np.array(img, dtype=np.float32))
for i in data[index][1]:
    x1 = int(i["points"][0]['x'] * IMAGE_WIDTH)
    x2 = int(i["points"][1]['x'] * IMAGE_WIDTH)
    y1 = int(i["points"][0]['y'] * IMAGE_HEIGHT)
    y2 = int(i["points"][1]['y'] * IMAGE_HEIGHT)
```

```
masks[index][y1:y2, x1:x2] = 1
```

- ▼ Print the shape of X_train and mask array (1 mark)

```
X_train.shape
```

↳ (200, 224, 224, 3)

```
masks.shape
```

↳ (200, 224, 224)

- ▼ Print a sample image and image array

```
from matplotlib import pyplot
n = 10
print(X_train[n])
pyplot.imshow(X_train[n])
```

↳

```
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or
[[[-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]

 ...
 [-1. -1. -1. ]
 [-1. -1. -1. ]
 [-1. -1. -1. ]]

[[[-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]

 ...
 [-1. -1. -1. ]
 [-1. -1. -1. ]
 [-1. -1. -1. ]]

[[[-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]
 [-0.98431373 -0.98431373 -0.98431373]

 ...
 [-1. -1. -1. ]
 [-1. -1. -1. ]
 [-1. -1. -1. ]]

...
[[[-1. -1. -1. ]
 [-1. -1. -1. ]
 [-1. -1. -1. ]

 ...
 [-0.96862745 -0.96862745 -0.96862745]
 [-0.96078432 -0.96078432 -0.96078432]
 [-0.96078432 -0.96078432 -0.96078432]]]

[[[-1. -1. -1. ]
 [-1. -1. -1. ]
 [-1. -1. -1. ]

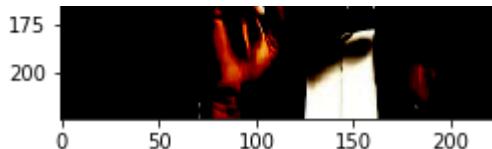
 ...
 [-0.96862745 -0.96862745 -0.96862745]
 [-0.96078432 -0.96078432 -0.96078432]
 [-0.95294118 -0.95294118 -0.95294118]]]

[[[-1. -1. -1. ]
 [-1. -1. -1. ]
 [-1. -1. -1. ]

 ...
 [-0.97647059 -0.97647059 -0.97647059]
 [-0.96862745 -0.96862745 -0.96862745]
 [-0.96078432 -0.96078432 -0.96078432]]]
```

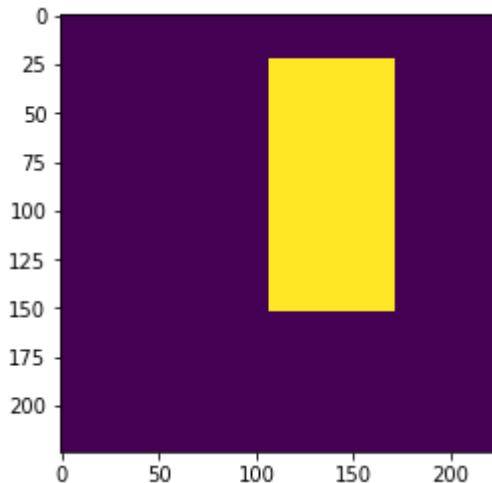
<matplotlib.image.AxesImage at 0x7feb01c7d908>





```
pyplot.imshow(masks[n])
```

↳ <matplotlib.image.AxesImage at 0x7feb01748898>



▼ Create the model (10 marks)

- Add MobileNet as model with below parameter values
 - input_shape: IMAGE_HEIGHT, IMAGE_WIDTH, 3
 - include_top: False
 - alpha: 1.0
 - weights: "imagenet"
- Add UNET architecture layers
 - This is the trickiest part of the project, you need to research and implement it correctly

```
from tensorflow.keras.applications.mobilenet import MobileNet
from tensorflow.keras.layers import Concatenate, UpSampling2D, Conv2D, Reshape
from tensorflow.keras.models import Model
```

```
def create_model(trainable=True):
    model = ##### Add your code here #####
    for layer in model.layers:
        layer.trainable = trainable

    # Add all the UNET layers here
    ##### Add your code here #####
```

```
return ##### Add your code here #####
```

```
class DataGenerator(Sequence):

    def __init__(self, csv_file):
        self.paths = []

        with open(csv_file, "r") as file:
            self.mask = np.zeros((sum(1 for line in file), HEIGHT_CELLS, WIDTH_CELLS))
            file.seek(0)

            reader = csv.reader(file, delimiter=",")

            for index, row in enumerate(reader):
                for i, r in enumerate(row[1:7]):
                    row[i+1] = int(r)

            path, image_height, image_width, x0, y0, x1, y1, _, _ = row

            path_mask = path.replace('.jpg','.png')

            mask_img = cv2.imread(path_mask)
            mask_img = (mask_img!=2)*1.0
            mask_img = cv2.resize(mask_img, (28, 28))
            mask_img = 1.0*(mask_img[:, :, 0]>0.2)
            self.mask[index, :, :] = np.squeeze(mask_img)

        self.paths.append(path)

    def __len__(self):
        return math.ceil(len(self.mask) / BATCH_SIZE)

    def __getitem__(self, idx):
        batch_paths = self.paths[idx * BATCH_SIZE:(idx + 1) * BATCH_SIZE]
        batch_masks = self.mask[idx * BATCH_SIZE:(idx + 1) * BATCH_SIZE]

        batch_images = np.zeros((len(batch_paths), IMAGE_HEIGHT, IMAGE_WIDTH, 3), dtype=np.float32)
        for i, f in enumerate(batch_paths):
            img = Image.open(f)
            img = img.resize((IMAGE_WIDTH, IMAGE_HEIGHT))
            img = img.convert('RGB')

            batch_images[i] = preprocess_input(np.array(img, dtype=np.float32))
            img.close()

        return batch_images, batch_masks
```

▼ Call the create_model function

```
# Give trainable=False as argument, if you want to freeze lower layers for fast training (
model = create_model()
```

```
# Print summary
model.summary()

def create_model(trainable=True):
    model = MobileNet(input_shape=(IMAGE_HEIGHT, IMAGE_WIDTH, 3), include_top=False, alpha
    for layer in model.layers:
        layer.trainable = trainable

    block1 = model.get_layer("conv_pw_5_relu").output
    block2 = model.get_layer("conv_pw_11_relu").output
    block3 = model.get_layer("conv_pw_13_relu").output

    x = Concatenate()([UpSampling2D()(block3), block2])
    x = Concatenate()([UpSampling2D()(x), block1])

    x = Conv2D(1, kernel_size=1, activation="sigmoid")(x)
    x = Reshape((HEIGHT_CELLS, WIDTH_CELLS))(x)

    return Model(inputs=model.input, outputs=x)
```

▼ Define dice coefficient function (5 marks)

- Create a function to calculate dice coefficient

```
def dice_coefficient(y_true, y_pred):
    numerator = 2 * tf.reduce_sum(y_true * y_pred)
    denominator = tf.reduce_sum(y_true + y_pred)

    return numerator / (denominator + tf.keras.backend.epsilon())
```

▼ Define loss

```
from tensorflow.keras.losses import binary_crossentropy
from tensorflow.keras.backend import log, epsilon
def loss(y_true, y_pred):
    return binary_crossentropy(y_true, y_pred) - log(dice_coefficient(y_true, y_pred)) + epsilon()

def loss(y_true, y_pred):
    return binary_crossentropy(y_true, y_pred) - tf.log(dice_coefficient(y_true, y_pred)) + epsilon()
```

▼ Compile the model (2 marks)

- Complie the model using below parameters
 - loss: use the loss function defined above
 - optimizers: use Adam optimizer

- metrics: use dice_coefficient function defined above

```
model = create_model(False)
model.summary()

train_datagen = DataGenerator(TRAIN_CSV)
validation_datagen = DataGenerator(VALIDATION_CSV)

optimizer = Adam(lr=1e-4, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgrad=False)
model.compile(loss=loss, optimizer=optimizer, metrics=[dice_coefficient])

checkpoint = ModelCheckpoint("model-{val_loss:.2f}.h5", monitor="val_loss", verbose=1, save_weights_only=True, mode="auto", period=1)
stop = EarlyStopping(monitor="val_loss", patience=PATIENCE, mode="auto")
reduce_lr = ReduceLROnPlateau(monitor="val_loss", factor=0.2, patience=5, min_lr=1e-6, ver
```

⇨

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python
Instructions for updating:
If using Keras pass *_constraint arguments to layers.
Downloading data from <https://github.com/fchollet/deep-learning-models/releases/download/17227776/17225924> [=====] - 1s 0us/step
Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[None, 224, 224, 3]	0	
conv1_pad (ZeroPadding2D)	(None, 225, 225, 3)	0	input_1[0][0]
conv1 (Conv2D)	(None, 112, 112, 32)	864	conv1_pad[0][0]
conv1_bn (BatchNormalization)	(None, 112, 112, 32)	128	conv1[0][0]
conv1_relu (ReLU)	(None, 112, 112, 32)	0	conv1_bn[0][0]
conv_dw_1 (DepthwiseConv2D)	(None, 112, 112, 32)	288	conv1_relu[0][0]
conv_dw_1_bn (BatchNormalization)	(None, 112, 112, 32)	128	conv_dw_1[0][0]
conv_dw_1_relu (ReLU)	(None, 112, 112, 32)	0	conv_dw_1_bn[0][0]
conv_pw_1 (Conv2D)	(None, 112, 112, 64)	2048	conv_dw_1_relu[0][0]
conv_pw_1_bn (BatchNormalization)	(None, 112, 112, 64)	256	conv_pw_1[0][0]
conv_pw_1_relu (ReLU)	(None, 112, 112, 64)	0	conv_pw_1_bn[0][0]
conv_pad_2 (ZeroPadding2D)	(None, 113, 113, 64)	0	conv_pw_1_relu[0][0]
conv_dw_2 (DepthwiseConv2D)	(None, 56, 56, 64)	576	conv_pad_2[0][0]
conv_dw_2_bn (BatchNormalization)	(None, 56, 56, 64)	256	conv_dw_2[0][0]
conv_dw_2_relu (ReLU)	(None, 56, 56, 64)	0	conv_dw_2_bn[0][0]
conv_pw_2 (Conv2D)	(None, 56, 56, 128)	8192	conv_dw_2_relu[0][0]
conv_pw_2_bn (BatchNormalization)	(None, 56, 56, 128)	512	conv_pw_2[0][0]
conv_pw_2_relu (ReLU)	(None, 56, 56, 128)	0	conv_pw_2_bn[0][0]
conv_dw_3 (DepthwiseConv2D)	(None, 56, 56, 128)	1152	conv_pw_2_relu[0][0]
conv_dw_3_bn (BatchNormalization)	(None, 56, 56, 128)	512	conv_dw_3[0][0]
conv_dw_3_relu (ReLU)	(None, 56, 56, 128)	0	conv_dw_3_bn[0][0]
conv_pw_3 (Conv2D)	(None, 56, 56, 128)	16384	conv_dw_3_relu[0][0]
conv_pw_3_bn (BatchNormalization)	(None, 56, 56, 128)	512	conv_pw_3[0][0]
conv_pw_3_relu (ReLU)	(None, 56, 56, 128)	0	conv_pw_3_bn[0][0]
conv_pad_4 (ZeroPadding2D)	(None, 57, 57, 128)	0	conv_pw_3_relu[0][0]
conv_dw_4 (DepthwiseConv2D)	(None, 28, 28, 128)	1152	conv_pad_4[0][0]

conv_dw_4_bn (BatchNormalization)	(None, 28, 28, 128)	512	conv_dw_4[0][0]
conv_dw_4_relu (ReLU)	(None, 28, 28, 128)	0	conv_dw_4_bn[0][0]
conv_pw_4 (Conv2D)	(None, 28, 28, 256)	32768	conv_dw_4_relu[0][0]
conv_pw_4_bn (BatchNormalization)	(None, 28, 28, 256)	1024	conv_pw_4[0][0]
conv_pw_4_relu (ReLU)	(None, 28, 28, 256)	0	conv_pw_4_bn[0][0]
conv_dw_5 (DepthwiseConv2D)	(None, 28, 28, 256)	2304	conv_pw_4_relu[0][0]
conv_dw_5_bn (BatchNormalization)	(None, 28, 28, 256)	1024	conv_dw_5[0][0]
conv_dw_5_relu (ReLU)	(None, 28, 28, 256)	0	conv_dw_5_bn[0][0]
conv_pw_5 (Conv2D)	(None, 28, 28, 256)	65536	conv_dw_5_relu[0][0]
conv_pw_5_bn (BatchNormalization)	(None, 28, 28, 256)	1024	conv_pw_5[0][0]
conv_pw_5_relu (ReLU)	(None, 28, 28, 256)	0	conv_pw_5_bn[0][0]
conv_pad_6 (ZeroPadding2D)	(None, 29, 29, 256)	0	conv_pw_5_relu[0][0]
conv_dw_6 (DepthwiseConv2D)	(None, 14, 14, 256)	2304	conv_pad_6[0][0]
conv_dw_6_bn (BatchNormalization)	(None, 14, 14, 256)	1024	conv_dw_6[0][0]
conv_dw_6_relu (ReLU)	(None, 14, 14, 256)	0	conv_dw_6_bn[0][0]
conv_pw_6 (Conv2D)	(None, 14, 14, 512)	131072	conv_dw_6_relu[0][0]
conv_pw_6_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_pw_6[0][0]
conv_pw_6_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_6_bn[0][0]
conv_dw_7 (DepthwiseConv2D)	(None, 14, 14, 512)	4608	conv_pw_6_relu[0][0]
conv_dw_7_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_dw_7[0][0]
conv_dw_7_relu (ReLU)	(None, 14, 14, 512)	0	conv_dw_7_bn[0][0]
conv_pw_7 (Conv2D)	(None, 14, 14, 512)	262144	conv_dw_7_relu[0][0]
conv_pw_7_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_pw_7[0][0]
conv_pw_7_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_7_bn[0][0]
conv_dw_8 (DepthwiseConv2D)	(None, 14, 14, 512)	4608	conv_pw_7_relu[0][0]
conv_dw_8_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_dw_8[0][0]
conv_dw_8_relu (ReLU)	(None, 14, 14, 512)	0	conv_dw_8_bn[0][0]
conv_pw_8 (Conv2D)	(None, 14, 14, 512)	262144	conv_dw_8_relu[0][0]
conv_pw_8_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_pw_8[0][0]
conv_pw_8_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_8_bn[0][0]
conv_dw_9 (DepthwiseConv2D)	(None, 14, 14, 512)	4608	conv_pw_8_relu[0][0]

conv_dw_9_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_dw_9[0][0]
conv_dw_9_relu (ReLU)	(None, 14, 14, 512)	0	conv_dw_9_bn[0][0]
conv_pw_9 (Conv2D)	(None, 14, 14, 512)	262144	conv_dw_9_relu[0][0]
conv_pw_9_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_pw_9[0][0]
conv_pw_9_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_9_bn[0][0]
conv_dw_10 (DepthwiseConv2D)	(None, 14, 14, 512)	4608	conv_pw_9_relu[0][0]
conv_dw_10_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_dw_10[0][0]
conv_dw_10_relu (ReLU)	(None, 14, 14, 512)	0	conv_dw_10_bn[0][0]
conv_pw_10 (Conv2D)	(None, 14, 14, 512)	262144	conv_dw_10_relu[0][0]
conv_pw_10_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_pw_10[0][0]
conv_pw_10_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_10_bn[0][0]
conv_dw_11 (DepthwiseConv2D)	(None, 14, 14, 512)	4608	conv_pw_10_relu[0][0]
conv_dw_11_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_dw_11[0][0]
conv_dw_11_relu (ReLU)	(None, 14, 14, 512)	0	conv_dw_11_bn[0][0]
conv_pw_11 (Conv2D)	(None, 14, 14, 512)	262144	conv_dw_11_relu[0][0]
conv_pw_11_bn (BatchNormalization)	(None, 14, 14, 512)	2048	conv_pw_11[0][0]
conv_pw_11_relu (ReLU)	(None, 14, 14, 512)	0	conv_pw_11_bn[0][0]
conv_pad_12 (ZeroPadding2D)	(None, 15, 15, 512)	0	conv_pw_11_relu[0][0]
conv_dw_12 (DepthwiseConv2D)	(None, 7, 7, 512)	4608	conv_pad_12[0][0]
conv_dw_12_bn (BatchNormalization)	(None, 7, 7, 512)	2048	conv_dw_12[0][0]
conv_dw_12_relu (ReLU)	(None, 7, 7, 512)	0	conv_dw_12_bn[0][0]
conv_pw_12 (Conv2D)	(None, 7, 7, 1024)	524288	conv_dw_12_relu[0][0]
conv_pw_12_bn (BatchNormalization)	(None, 7, 7, 1024)	4096	conv_pw_12[0][0]
conv_pw_12_relu (ReLU)	(None, 7, 7, 1024)	0	conv_pw_12_bn[0][0]
conv_dw_13 (DepthwiseConv2D)	(None, 7, 7, 1024)	9216	conv_pw_12_relu[0][0]
conv_dw_13_bn (BatchNormalization)	(None, 7, 7, 1024)	4096	conv_dw_13[0][0]
conv_dw_13_relu (ReLU)	(None, 7, 7, 1024)	0	conv_dw_13_bn[0][0]
conv_pw_13 (Conv2D)	(None, 7, 7, 1024)	1048576	conv_dw_13_relu[0][0]
conv_pw_13_bn (BatchNormalization)	(None, 7, 7, 1024)	4096	conv_pw_13[0][0]
conv_pw_13_relu (ReLU)	(None, 7, 7, 1024)	0	conv_pw_13_bn[0][0]
up_sampling2d (UpSampling2D)	(None, 14, 14, 1024)	0	conv_pw_13_relu[0][0]

concatenate (Concatenate)	(None, 14, 14, 1536) 0	up_sampling2d[0][0] conv_pw_11_relu[0][0]
up_sampling2d_1 (UpSampling2D)	(None, 28, 28, 1536) 0	concatenate[0][0]
concatenate_1 (Concatenate)	(None, 28, 28, 1792) 0	up_sampling2d_1[0][0] conv_pw_5_relu[0][0]
conv2d (Conv2D)	(None, 28, 28, 1) 1793	concatenate_1[0][0]
reshape (Reshape)	(None, 28, 28) 0	conv2d[0][0]
<hr/>		
Total params:	3,230,657	
Trainable params:	1,793	
Non-trainable params:	3,228,864	

```

-----  

FileNotFoundError                                     Traceback (most recent call last)  

<ipython-input-30-88545abbee4c> in <module>()  

      2 model.summary()  

      3  

----> 4 train_datagen = DataGenerator(TRAIN_CSV)  

      5 validation_datagen = DataGenerator(VALIDATION_CSV)  

      6  

<ipython-input-26-57ac1072ace7> in __init__(self, csv_file)  

      4         self.paths = []  

      5  

----> 6     with open(csv_file, "r") as file:  

      7         self.mask = np.zeros((sum(1 for line in file), HEIGHT_CELLS, WIDTH_CELLS))  

      8         file.seek(0)

```

FileNotFoundError: [Errno 2] No such file or directory: '/content/drive/My Drive/Cola

SEARCH STACK OVERFLOW

▼ Define checkpoint and earlystopping

```

from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
checkpoint = ModelCheckpoint("model-{loss:.2f}.h5", monitor="loss", verbose=1, save_best_only=True)
https://colab.research.google.com/drive/1JDija1itRxz4m62DzbuFSuwWRmdbX3VY#scrollTo=Z0Qa1UOue9TE&printMode=true

```

```

        save_weights_only=True, mode="min", period=1)
stop = EarlyStopping(monitor="loss", patience=5, mode="min")
reduce_lr = ReduceLROnPlateau(monitor="loss", factor=0.2, patience=5, min_lr=1e-6, verbose
    ↴ WARNING:tensorflow:`period` argument is deprecated. Please use `save_freq` to specify

```

▼ Fit the model (2 marks)

- Fit the model using below parameters
 - epochs: you can decide
 - batch_size: 1
 - callbacks: checkpoint, reduce_lr, stop

```

model.fit_generator(generator=train_datagen,
                    epochs=EPOCHS,
                    validation_data=validation_datagen,
                    callbacks=[checkpoint, reduce_lr, stop],
                    workers=THREADS,
                    use_multiprocessing=False,
                    shuffle=True,
                    verbose=1)

```

▼ Get the predicted mask for a sample image (3 marks)

```

n = 10
sample_image = X_train[n]

import cv2
import numpy as np

from keras.applications.mobilenet import preprocess_input
    ↴ Using TensorFlow backend.

```

▼ Impose the mask on the image (3 marks)

```

WEIGHTS_FILE = "model-0.92.h5"
THRESHOLD = 0.8
EPSILON = 0.02

model = create_model()
model.load_weights(WEIGHTS_FILE)

```

