

▼ Build a DNN using Keras with RELU and ADAM

▼ Load tensorflow

```
import tensorflow as tf  
tf.set_random_seed(42)
```

☞ The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.
We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %tenc

▼ Collect Fashion mnist data from tf.keras.datasets

```
(trainX, trainY), (testX, testY) = tf.keras.datasets.mnist.load_data()
```

☞ Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist/>.
11493376/11490434 [=====] - 0s 0us/step

▼ Change train and test labels into one-hot vectors

```
trainY = tf.keras.utils.to_categorical(trainY, num_classes=10)  
testY = tf.keras.utils.to_categorical(testY, num_classes=10)
```

Build the Graph

▼ Initialize model, reshape & normalize data

```
#Initialize model, reshape & normalize data  
tf.keras.backend.clear_session()  
model = tf.keras.models.Sequential()  
model.add(tf.keras.layers.Reshape((784,), input_shape=(28, 28,)))  
model.add(tf.keras.layers.BatchNormalization())
```

▼ Add two fully connected layers with 200 and 100 neurons respectively with relu activation

```
#Hidden layers  
model.add(tf.keras.layers.Dense(200, activation='relu', name='Layer_1'))  
  
model.add(tf.keras.layers.BatchNormalization())  
  
model.add(tf.keras.layers.Dense(100, activation='relu', name='Layer_2'))
```

```
model.add(tf.keras.layers.BatchNormalization())
#Dropout layer
model.add(tf.keras.layers.Dropout(0.25))
```

Add the output layer with a fully connected layer with 10 neurons with softmax activation, categorical_crossentropy loss and adam optimizer and train the network. And, run the model.

```
#Output layer
model.add(tf.keras.layers.Dense(10, activation='softmax', name='Output'))

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()
```

↳ Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
reshape (Reshape)	(None, 784)	0
batch_normalization (BatchNormalizer)	(None, 784)	3136
Layer_1 (Dense)	(None, 200)	157000
batch_normalization_1 (BatchNormalizer)	(None, 200)	800
Layer_2 (Dense)	(None, 100)	20100
batch_normalization_2 (BatchNormalizer)	(None, 100)	400
dropout (Dropout)	(None, 100)	0
Output (Dense)	(None, 10)	1010
<hr/>		
Total params: 182,446		
Trainable params: 180,278		
Non-trainable params: 2,168		

```
model.fit(trainX,trainY,
           validation_data=(testX,testY),
           epochs=5,
           batch_size=32)
```

↳

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/5
60000/60000 [=====] - 12s 197us/sample - loss: 0.1831 - acc: 0.
Epoch 2/5
60000/60000 [=====] - 11s 189us/sample - loss: 0.1282 - acc: 0.
Epoch 3/5
60000/60000 [=====] - 11s 187us/sample - loss: 0.1003 - acc: 0.
Epoch 4/5
60000/60000 [=====] - 11s 191us/sample - loss: 0.0896 - acc: 0.
Epoch 5/5
60000/60000 [=====] - 11s 188us/sample - loss: 0.0745 - acc: 0.
<tensorflow.python.keras.callbacks.History at 0x7f4cea3665c0>
```